

Check that internal equilibrium solutions are identical

This package contains two sets of code for returning the equilibrium solution to the Griffin et al. 2014 transmission model. The first was written in compiled R code by Jamie Griffin, and the second was written in Rcpp by Xiaoyu Li, before being adapted into the malariaModelFit package by Bob Verity. Here we will compare both sets of code to establish 1) that both sets of code produce the same results, 2) what speed gain we achieve by moving the compiled R function to Rcpp.

Compare solutions

Required inputs when determining the equilibrium solution include; entomological inoculation rate (EIR), treatment coverage (ft), a vector of parameters (p), the age distribution of the population (age), and Gaussian quadrature parameters capturing biting heterogeneity (h). We will fix age and biting heterogeneity and draw `n` random values for the other inputs.

```
## Loading malariaModelFit

## Warning in setup_ns_exports(pkg, export_all): Objects listed as exports,
## but not present in namespace: human_equilibrium_fast, plot_nice1

n <- 1e3 # number of random inputs to explore

age <- default_age()
h <- gq_normal(9)
EIR <- runif(n,1,100)
ft <- runif(n)
p <- random_parameters(n)
```

Now calculate equilibrium solution using both functions, and check that results are equal using `all.equal`, which checks for “near equality”.

```
is_equal <- rep(NA,n)
for (i in 1:n) {
  eq1 <- human_equilibrium(EIR=EIR[i], ft=ft[i], p=p[[i]], age=age, h=h)
  eq2 <- human_equilibrium_cpp(EIR=EIR[i], ft=ft[i], p=p[[i]], age=age, h$nodes, h$weights)
  is_equal[i] <- all.equal(eq1,eq2)
}

# check that results match
if (all(is_equal)) {
  cat("Good news - solutions match exactly!\n")
} else {
  warning("Solutions do not match")
}
```

```
## Good news - solutions match exactly!
```

Hopefully you got good news when running the above, otherwise check your code for bugs!

Compare speed

We will compare speed with the `microbenchmark` package. Here we compare evaluation time for each of the `n` parameter sets generated above, repeating the evaluation `repeats` times in each case. Results are stored in the list `res`, and the median over repeats is stored in the matrix `res_median`.

```

library(microbenchmark)

repeats <- 1 # number of times to repeat evaluation for each parameter set

# loop over parameter sets
res <- NULL
res_median <- matrix(NA,n,2)
for (i in 1:n) {

  # benchmark functions
  res_i <- microbenchmark(
    human_equilibrium(EIR=EIR[i], ft=ft[i], p=p[[i]], age=age, h=h),
    human_equilibrium_cpp(EIR=EIR[i], ft=ft[i], p=p[[i]], age=age, h$nodes, h$weights),
    times=repeats
  )

  # store results
  res$expr <- c(res$expr, res_i$expr)
  res$time <- c(res$time, res_i$time)
  res_median[i,] <- summary(res_i)$median
}


```

We can now produce a box-plot over all results:

```

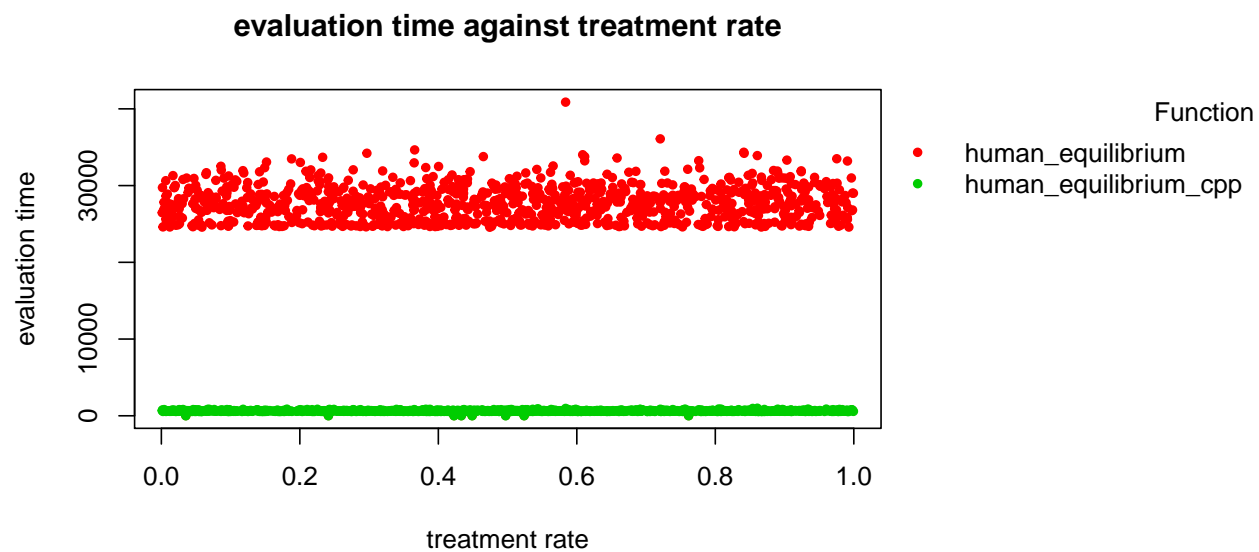
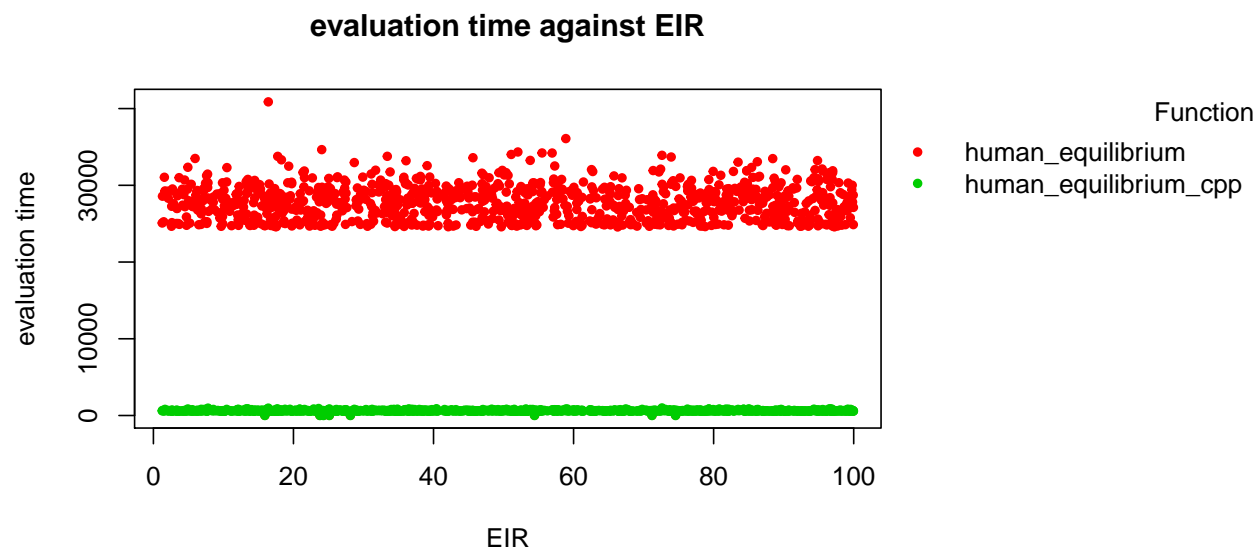
functionNames <- c("human_equilibrium","human_equilibrium_cpp")
boxplot(split(res$time, f=res$expr), names=functionNames, ylab="evaluation time")

```



Result: In this evaluation human_equilibrium_cpp was 43.35 times faster than human_equilibrium.

We can also dig deeper into timings by plotting evaluation time against EIR or treatment rate:



Neither plot shows any strong relationship with evaluation time.

(Rmarkdown last evaluated: 2018-01-12 11:07:01)